

6. FinOps Analysis of Results

6.1 - Scope and Method for FinOps KPI Analysis

The analytical question is whether the published products make cost and usage information interpretable for FinOps decisions. Building a dashboard is insufficient evidence by itself: a displayed figure must be tied to a named measure, grain, period, filter, source table and reconciliation rule. This chapter therefore uses the same definition for an indicator when it appears in several views and checks that subject-oriented datamarts agree on shared totals. Section 7.1 reports the local technical checks; this chapter examines what their values can and cannot mean to a business user.

The quantified dashboard observations currently documented come from January 2025 in the frozen DuckDB POC: 164,145 charge lines, EUR 912,000 in billed cost and fourteen published datamarts, eleven consumed by Streamlit. Section 7.1 presents the local validation scope. The generator has a longer synthetic Daily history, and the Belgium working note reports an eighteen-month DEV/PROD billing backfill. Those facts define potential analysis coverage, but local January indicators cannot be combined with cloud monthly indicators until the cloud outputs, executed revision and reconciliation evidence are recorded. An eighteen-month synthetic trend would illustrate the analytical workflow, not the real evolution of the host organization's expenditure.

For each KPI, the analysis records its business purpose, FOCUS input fields, formula, unit, charge-line or aggregate grain, period, exclusions, source datamart and interpretation. `BilledCost` describes the billed amount represented by the source; `ListCost`, `ContractedCost` and `EffectiveCost` express different price perspectives. Arithmetic differences among them are comparisons, not automatically realized savings. In particular, eligibility, charge categories, credits, commitments and scope must be reviewed before labeling a difference as an optimization result.

The dashboard subjects are executive overview, price comparisons, allocation, resource drivers and data quality. An overview should reconcile to the same billed total as the monthly billing datamart. Allocation slices should reconcile to that total after explicit treatment of unknown or unassigned values. Resource rankings identify where to investigate; without utilization, criticality and business context, high cost alone does not identify waste. Data-quality indicators and audit freshness determine whether a business user should trust a period's analysis. The operational and architecture pages explain lineage and status but are not substitutes for validated cost KPIs.

The local POC supports an Inform demonstration and structured investigation of possible Optimize questions. It has not established realized company savings, automated remediation,

identity-based filtering or sustained Operate outcomes. The distinction between demonstrated behavior and intended use governs the interpretations in sections 6.2 and 6.3.

6.2 - Analysis of Costs and Pricing Mechanisms

The executive datamart reports EUR 912,000.00 in `BilledCost` for 164,145 charge lines. The corresponding `ListCost` is EUR 1,290,989.95, `ContractedCost` EUR 1,044,652.85 and `EffectiveCost` EUR 1,101,755.75. Section 7.1 summarizes this local POC validation. Billed and effective cost have different FOCUS meanings and therefore should not be expected to coincide simply because both appear in a dashboard. Their divergence requires interpretation against charge categories and the synthetic generator's transformations before it can inform a financial decision.

The list-to-effective difference is EUR 189,234.20, or 14.66% of list cost. This describes a price-metric gap within the synthetic dataset. It is not a saving produced by the platform, nor evidence that the organization could recover this amount. The contracted-to-effective difference is negative EUR 57,102.90. Describing it as positive commitment savings would reverse the observed sign and misrepresent the data. Before using either gap for optimization, the analysis must align charge eligibility, credits, negotiated rates, commitment-covered usage, amortization and aggregation grain. In the present dataset, the safe label is “cost comparison,” not “realized savings.”

The same local datamart allows cost-driver analysis. Virtual Machines account for approximately EUR 195,399.16 and Storage Accounts for EUR 176,985.27. Together, their unrounded source values represent EUR 372,384.42, or 40.83% of the month's billed total. This concentration gives an analyst a sensible starting point for investigation, but does not establish that these services are overprovisioned or that their cost can be reduced. Utilization, commitments, workload criticality and owner context would be needed for an actionable recommendation.

These results demonstrate that the published subjects can reconcile common measures and expose nontrivial economic questions. They do not validate a company-specific rate policy, a cross-month trend or a monetary benefit. Extending the analysis to the eighteen cloud billing months requires datamart outputs and their Databricks run evidence, rather than extrapolation from January's local POC.

6.3 - Allocation and Resource Analysis

The allocation datamart groups January cost by cost center and service. Its groups reconcile to the EUR 912,000 billed total. However, EUR 875,320.02, or 95.98%, falls under cost center `Unknown`. Only EUR 36,679.98 is assigned to the three named synthetic centers in this view. The result shows limited attribution coverage: a center-level showback can display the data, but it cannot fairly allocate most of the amount to named budget owners. Improving the tagging or mapping rule would have greater immediate value than refining the chart used to present these groups.

Application ownership has a different coverage pattern. Rows with a missing or placeholder `application_owner_id` represent EUR 146,649.15, or 16.08% of the same total. The cost-center and application-owner fields describe different organizational dimensions, so these percentages are not interchangeable or additive. An application owner can be present while the cost center is unknown. Neither field proves that a person has accepted budget responsibility; ownership definitions and access rights still require business validation.

The local executive summary counts 22,453 active resources and 59 consumed services. The most expensive services are Virtual Machines and Storage Accounts, which together account for 40.83% of billed cost as discussed in section 6.2. Resource rankings can help prioritize investigation, but a high charge does not by itself indicate waste. A resource's function, utilization, resilience requirement and commitment coverage must be checked before suggesting a change.

These figures show why data quality and allocation are part of FinOps analysis rather than only pipeline checks. A technically reconciled EUR 912,000 total can still be weak evidence for accountability when most cost-center values are unknown. The finding is specific to the generated January dataset and must not be presented as a diagnosis of Technip Energies' actual tagging or allocation practice.